

# Introduction to xv6

By Ghazal Sadeghian

# What is xv6 ?

- simplified and lightweight OS based on Sixth Edition Unix
- It was designed in MIT for educational purposes
- theoretically actually boots on real 32-bit x86 hardware(and supports multicore)

(but we'll run it only single-core, in qemu emulator)

- It has an easy-to-understand structure

# How to run xv6?

XV6 can be installed as a standalone OS (not recommended)

XV6 can run in an emulated environment (recommended)

***QEMU*** is the emulator we use to run XV6

(QEMU is a user program/system emulator)

## xv6 technical requirements

- Linux environment
- Qemu simulator

---

# Getting xv6 running

**Step 1 – Install qemu:** [\\$ sudo apt install qemu](#)

**Step 2 – Install xv6:** Create a directory, and clone xv6 to that directory:

[\\$ git clone git://github.com/mit-pdos/xv6-public.git](#)

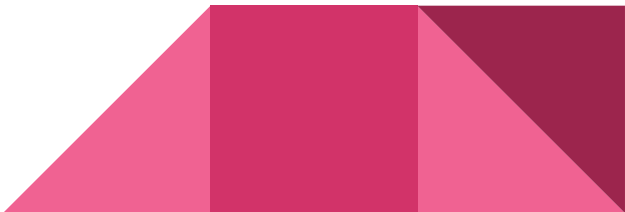
**Step 3 – Compile xv6 :** In the newly created xv6 directory: [\\$ make](#)

**Step 4 – Compile and run the emulator qemu:** [\\$ make qemu](#)

[\\$ make qemu-nox](#)


`make qemu` : runs the graphical version of qemu

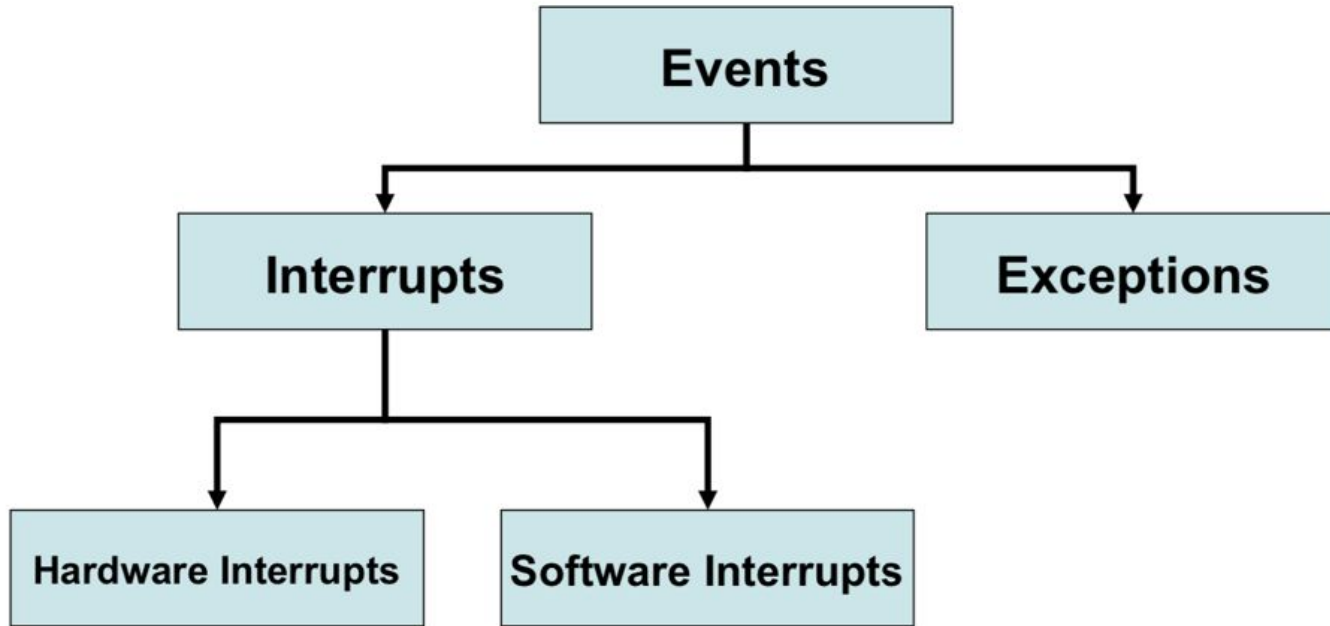
`make qemu-nox`: runs the console version



# Why event driven design?


- OS cannot trust user processes
  - User processes may be buggy or malicious
  - User process crash should not affect OS
- OS needs to guarantee fairness to all user processes
  - One process cannot 'hog' CPU time
  - Timer interrupts





# Hardware vs Software Interrupt

## Hardware Interrupt



A device (like PIC)

asserts a pin in the CPU

## Software Interrupt

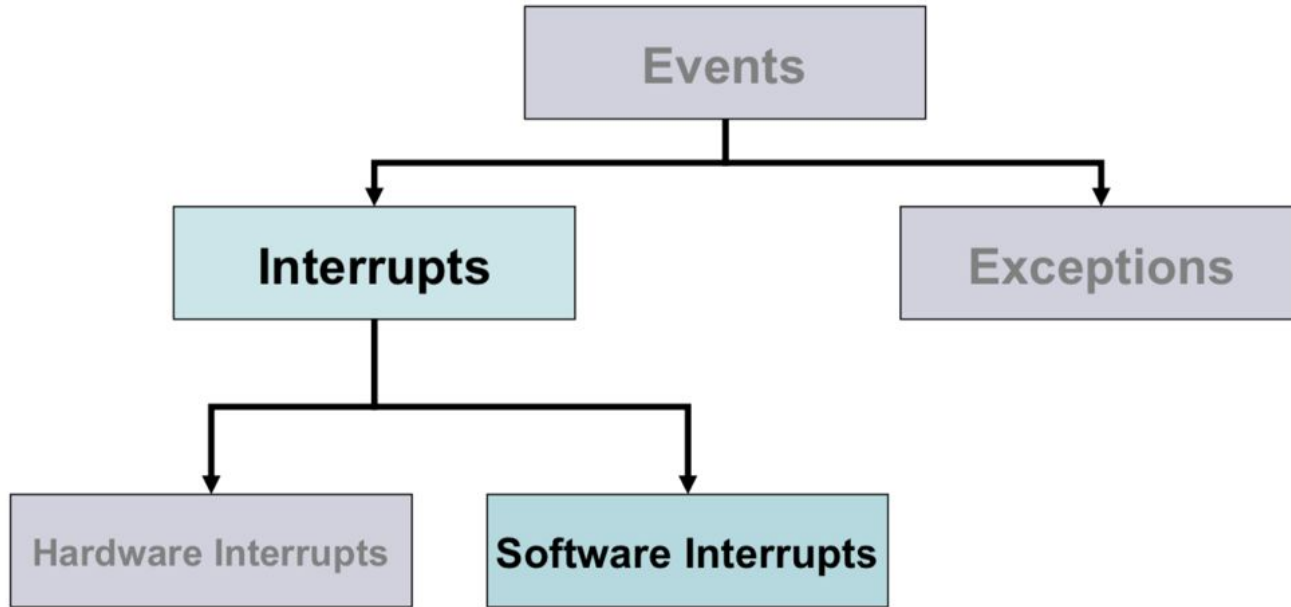


An instruction which when executed

causes an interrupt



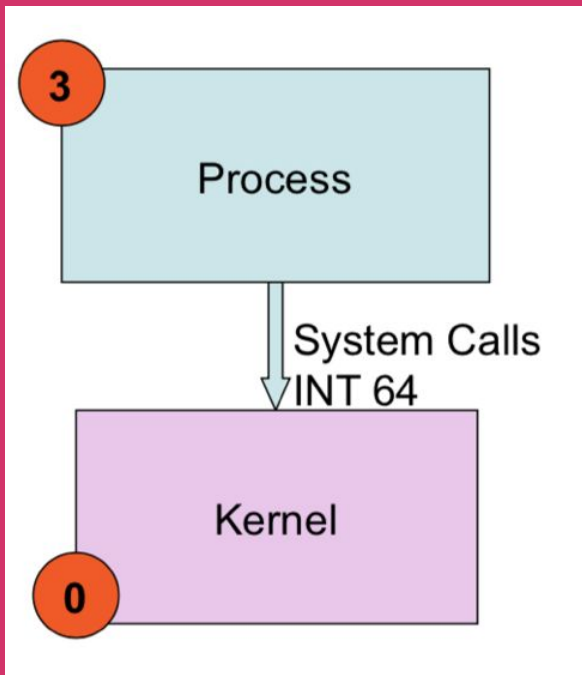
# System Calls



Software interrupt used for implementing system calls

- In Linux INT 128, is used for system calls

- In xv6, INT 64 is used for system calls



# System calls in xv6

- fork, exec, exit, wait, kill, getpid — process control
- open, read, write, close, fstat, dup — file operations
- mknod, unlink, link, chdir — directory operations

System call	Description
fork()	Create process
exit()	Terminate current process
wait()	Wait for a child process to exit
kill(pid)	Terminate process pid
getpid()	Return current process's id
sleep(n)	Sleep for n seconds
exec(filename, *argv)	Load a file and execute it
sbrk(n)	Grow process's memory by n bytes
open(filename, flags)	Open a file; flags indicate read/write
read(fd, buf, n)	Read n bytes from an open file into buf
write(fd, buf, n)	Write n bytes to an open file
close(fd)	Release open file fd
dup(fd)	Duplicate fd
pipe(p)	Create a pipe and return fd's in p
chdir(dirname)	Change the current directory
mkdir(dirname)	Create a new directory
mknod(name, major, minor)	Create a device file
fstat(fd)	Return info about an open file
link(f1, f2)	Create another name (f2) for the file f1
unlink(filename)	Remove a file

**How does the  
OS distinguish  
between the system  
calls?**

System  
call number

mov x, %eax  
INT 64

Based on the system call number  
function syscall invokes the  
corresponding syscall handler

## System call numbers

```
#define SYS_fork      1
#define SYS_exit      2
#define SYS_wait      3
#define SYS_pipe      4
#define SYS_read       5
#define SYS_kill       6
#define SYS_exec       7
#define SYS_fstat      8
#define SYS_chdir      9
#define SYS_dup      10
#define SYS_getpid     11
#define SYS_sbrk       12
#define SYS_sleep      13
#define SYS_uptime     14
#define SYS_open       15
#define SYS_write      16
#define SYS_mknod      17
#define SYS_unlink     18
#define SYS_link       19
#define SYS_mkdir      20
#define SYS_close      21
```

## System call handlers

```
[SYS_fork]    sys_fork,
[SYS_exit]    sys_exit,
[SYS_wait]    sys_wait,
[SYS_pipe]    sys_pipe,
[SYS_read]    sys_read,
[SYS_kill]    sys_kill,
[SYS_exec]    sys_exec,
[SYS_fstat]   sys_fstat,
[SYS_chdir]   sys_chdir,
[SYS_dup]     sys_dup,
[SYS_getpid]  sys_getpid,
[SYS_sbrk]    sys_sbrk,
[SYS_sleep]   sys_sleep,
[SYS_uptime]  sys_uptime,
[SYS_open]    sys_open,
[SYS_write]   sys_write,
[SYS_mknod]   sys_mknod,
[SYS_unlink]  sys_unlink,
[SYS_link]    sys_link,
[SYS_mkdir]   sys_mkdir,
[SYS_close]   sys_close,
```

ref : syscall.h, syscall() in syscall.c

SYSCALL([NAME]) in usys.S:

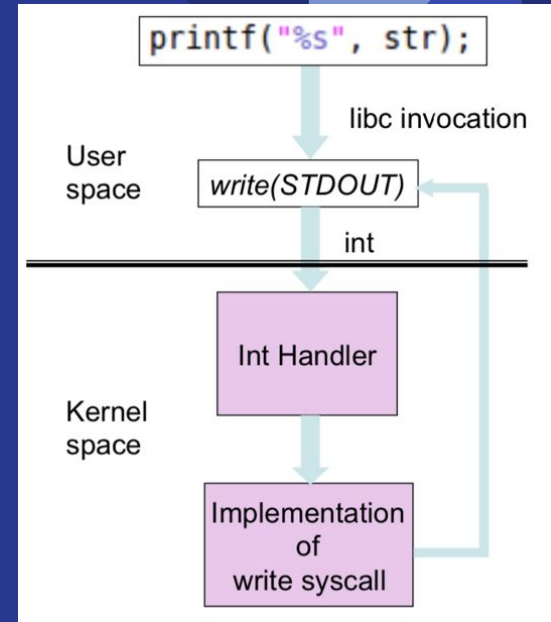
This line translates to the following assembly:

```
.globl name;
```

```
name:
```

```
mov $SYS_name, %eax; //putting the system call number in eax
```

```
int $0x40; //calling the system call handler in interrupt mode(64)
```



# Syscall(void)

All system calls are handled in this function.

The sys num which is saved in eax register is retrieved and system call is read from table.

```
void
syscall(void)
{
    int num;

    num = proc->tf->eax;
    if(num >= 0 && num < NELEM(syscalls) && syscalls[num])
        proc->tf->eax = syscalls[num]();
    else {
        cprintf("%d %s: unknown sys call %d\n",
                proc->pid, proc->name, num);
        proc->tf->eax = -1;
    }
}
```

# Prototype of a Typical System Call

```
int system_call( resource_descriptor, parameters)
```



return is generally  
'int' (or equivalent)  
sometimes 'void'

int used to denote completion  
status of system call sometimes  
also has additional information  
like number of bytes written to  
file

What OS resource is the target  
here?  
For example a file, device, etc.

If not specified, generally means  
the current process

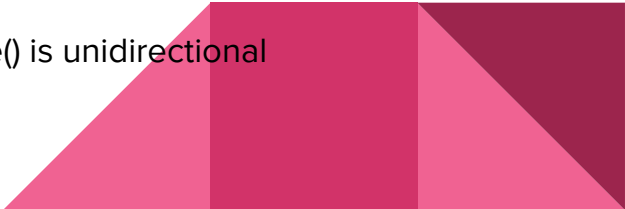
System call specific parameters  
passed.  
How are they passed?

# Question 1

## Fork()

- To create child process we use fork(). fork() returns :
  - **<0** fail to create child (new) process
  - **=0** Returned to the newly created child process
  - **>0** Returned to parent or caller. The value contains process ID of newly created child process..

## pipe()

- pipe() is used for passing information from one process to another. pipe() is unidirectional
- 



# getChildren

- getppid
- getchildren(pid)



# getCount

- number of times the referenced system call was invoked by the calling process

